

LAW4039 Notes

Adithya C. Ganesh

September 27, 2019

Contents

1	Lecture 1: Society, Institutions, Risk, and AI	2
1.1	R1: Why the US Went to War in Vietnam	2
1.2	R2: Five years after Deepwater Horizon, what has the disaster taught us?	2
1.3	R3: A Declaration of the Independence of Cyberspace (1996)	2
1.4	R4: The Great AI Awakening, N.Y. Times Magazine (Dec. 14, 2016)	2
1.5	R5: Regulating AI Systems: Risks, Challenges, Competencies, and Strategies	2
1.6	R6: Russell and Norvig, AI (2015)	2
1.7	R7: Suleyman and Musk: Humans Must Become Cyborgs to Avoid AI Domination	2
1.8	R8: Chui: Applying AI for Social Good, Discussion Paper, McKinsey and Co	2
1.9	Lecture Notes	2
1.9.1	A tragic intrusion	2

1 Lecture 1: Society, Institutions, Risk, and AI

1.1 R1: Why the US Went to War in Vietnam

<https://www.fpri.org/article/2017/04/united-states-went-war-vietnam/>

1.2 R2: Five years after Deepwater Horizon, what has the disaster taught us?

<https://www.minnpost.com/earth-journal/2015/04/five-years-after-deepwater-horizon-what-has-disaster-taught-us/>

1.3 R3: A Declaration of the Independence of Cyberspace (1996)

1.4 R4: The Great AI Awakening, N.Y. Times Magazine (Dec. 14, 2016)

1.5 R5: Regulating AI Systems: Risks, Challenges, Competencies, and Strategies

1.6 R6: Russell and Norvig, AI (2015)

1.7 R7: Suleyman and Musk: Humans Must Become Cyborgs to Avoid AI Domination

1.8 R8: Chui: Applying AI for Social Good, Discussion Paper, McKinsey and Co

1.9 Lecture Notes

The challenge for technologists: learning a little bit of law during these 10 weeks.

1.9.1 A tragic intrusion

The hardest cases to discuss and analyze: the death penalty cases (these ones go to California Supreme Court). Example: intruder enters the home of a woman who's home with her four year old son. Mother was stabbed, and bleeds to death.

Important questions:

- How to structure a six-picture "show-up" for the child at the push of a button?
- Do we use pervasive surveillance capacity into an ML-enabled infrastructure for security?

Death penalty cases are not-unlike analyzing the ethics of AI systems, where one must weigh the balance between different foundational ethical principles.

Two AI based approaches:

- GANs to generate images based on textual descriptions.
- Convolutional neural networks for facial recognition / person tracking.

But there are concerns:

- How accurate is the technology?
- What is the line between police and non-police?